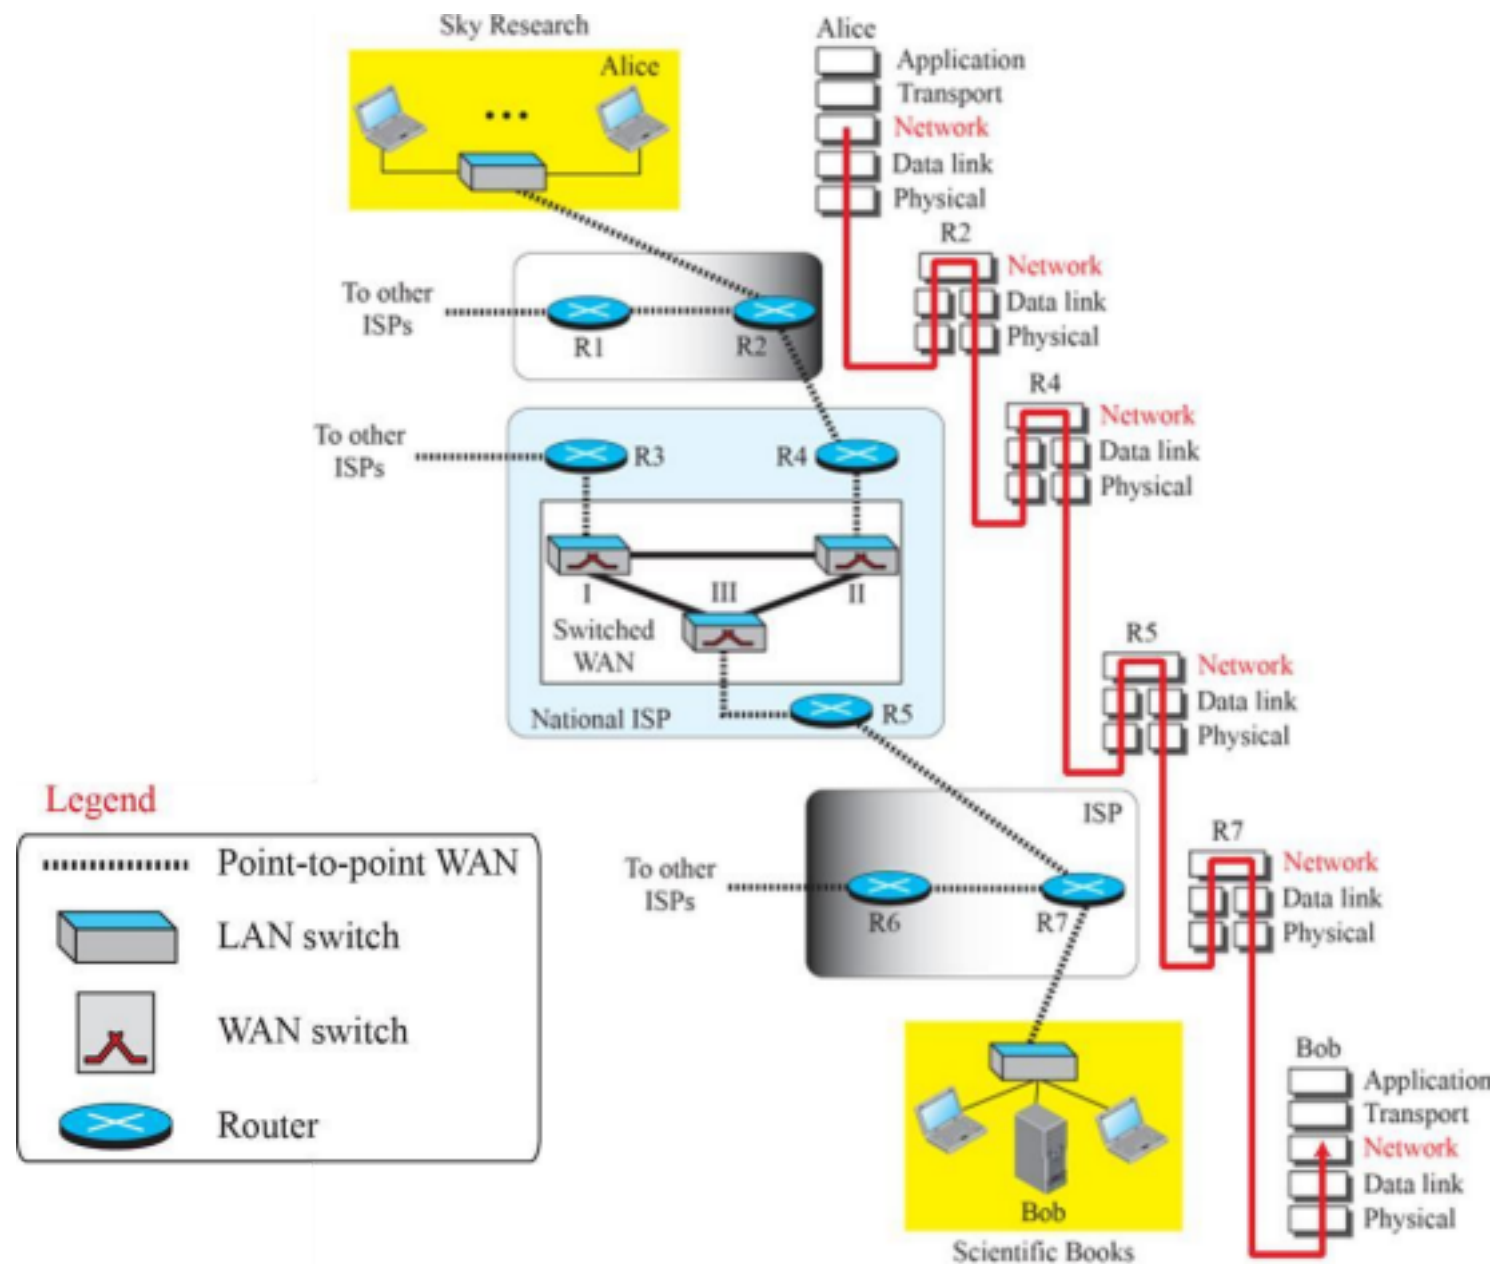


Chapter 4Networklayer

Objective

- We first discuss services that can be provided at the networklayer:packetizing, routing, and forwarding.
- We then discuss the network layer at the TCP/IP suite: IPv4andICMPv4. We also discuss IPv4 addressing and relatedissues.
- We then concentrate on the unicast routing and unicast routingprotocols.
- We then move to multicasting and multicast routing protocolsprotocol.



Network-Layer Services

Before discussing the network layer in the Internet today, let's briefly discuss the network-layer services that, in general, are expected from a network-layer protocol. Before

discussing the network layer in the Internet today, let's briefly discuss the network-layer services that, in general, are expected from a network-layer protocol.

- Packetizing
- Routing
- Forwarding

Packetizing

The network layer has an important job. It takes data from one place and sends it to another

without altering or using the data. It's like a postal service that delivers packages without changing their contents. Here's how it works:

- The source computer gets data from a higher-level program and puts it in a network packet. **This packet includes sender and receiver information** and other details needed for the network.
- The source computer then sends the packet to the data-link layer for delivery. **It can't change the data unless it's too big**, and it needs to split it into smaller parts.
- The destination computer gets the network packet from its data-link layer. It takes out the data and gives it to the right program. **If the packet was split into parts, the network layer waits until all parts arrive, puts them together**, and then gives the data to the program.
- Routers in between don't open the packets unless they need to

split them into smaller parts. They also don't change the sender and receiver info. They only look at the info to know where to send the packet next. If a packet is split, they copy the header to all parts and make some changes as needed.

Routing

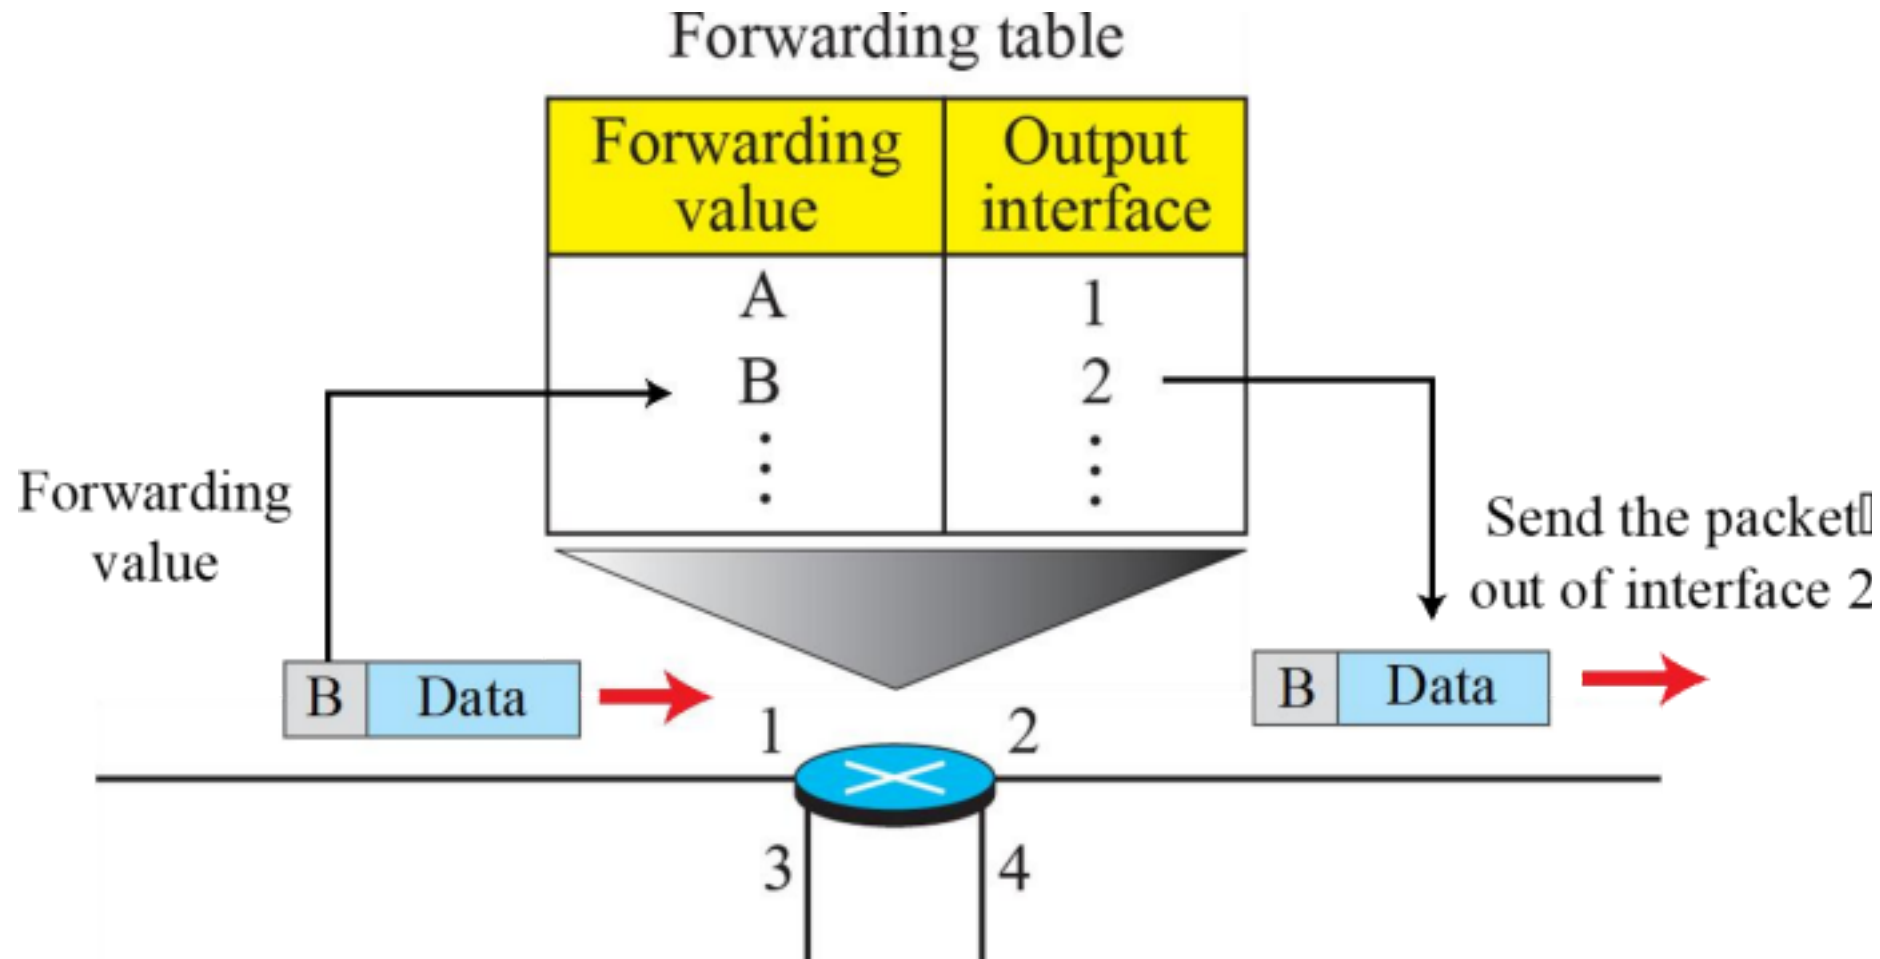
The network layer has another important job called "routing." This means it figures out **the best path for data to travel** from where it starts to where it needs to go. Think of it like finding the fastest way to get somewhere on a road trip. In a computer network, there are different paths data can take, like different roads on a map. The network layer's job is

topickthebest road for the data to travel on. To do this, the network layer has special plans, like GPS for data. **It uses something called "routing protocols" to help routers (like traffic cops for data) know which path to choose.** These plans are set up before any data starts moving.

Forwarding

Think of routing like making a plan and using special rules to decide how data should travel through a network. Now, **forwarding is what each router does when data actually arrives at its doorstep.** When data arrives at a router, it

needs to know **where to send it next**. Imagine it's like a traffic cop directing cars at an intersection. **The router looks at the data's address (or a special tag) and checks a list (sometimes called a forwarding table, or routing table) to figure out which way to send it.** So, routing is the plan-making part, and forwarding is the action part where the router decides where data goes based on that plan. It's like the router is following a map to guide the data to its destination.



Error control

In Chapter 3, we talked about how we handle mistakes in data transfer at the transport

layer. In Chapter 5, we'll dive into how we deal with errors at the data-link layer. Now, even though we can also handle errors at the network layer, **the folks who designed the network layer for the Internet didn't focus on it much.** One big reason is that

data in the network layer can get broken up into pieces at different points along its journey, like **splitting a puzzle into parts.** Checking for errors in this situation is not

very efficient. However, these designers **did add a special "checksum"**

field to the header of the data packet. It's like a security seal that makes sure the header isn't tampered with during its journey from one place to another. It's

important to note that even though the Internet's network layer doesn't directly fix errors, it has a helper **called ICMP. This protocol helps out if a data packet is thrown**

away or has weird stuff in its header. We'll get into more detail about ICMP later in this chapter.

Flow control

Flow control is like managing how much data a sender computer can send to a receiver computer without causing a traffic jam. If the sender's computer is making data faster than the receiver's computer can handle, it's like pouring too much water into a glass—it overflows. To avoid this, the receiver needs to let the sender know when it's getting overwhelmed with data. Interestingly, the network layer in the Internet doesn't directly handle this flow control job. Instead, it lets the sender send data whenever it's ready, without worrying about whether the receiver can keep up. There are a few reasons for this approach.

First, **the receiver's job at the network layer is pretty simple** because it doesn't have to deal with

error checking, so it's less likely to get overwhelmed. Second, the higher-level programs using the network layer can **manage the incoming data** by storing it temporarily and don't need to consume it all at once. Third, most of these higher-level programs **have their own flow control**, so adding more control at the network layer would make things more complex and less efficient.

Congestion control

Another thing to consider in network-layer protocols is congestion control. Congestion happens when there are too many data packets in one part of the Internet. It's like a traffic jam on the information highway. This can occur if computers send more data than the network or routers can handle. When this occurs, some packets might get dropped by routers to ease the traffic.

But here's the tricky part: **if too many packets get dropped, it can make things worse. That's because higher-level error checks might make the sender resend the lost packets, causing even more congestion. In extreme cases, everything**

can grind to a halt, and no data can get through.

We'll go into more detail about how to deal with congestion in the network layer later in this chapter, even though the Internet doesn't directly handle it.

Quality of service

With the Internet, we can now do cool things like video chats and streaming music. But for these things to work smoothly, we need to make sure the quality of our communication is really good. The Internet has gotten better at this by improving the quality of service (QoS), which is like making sure your online experience is top-notch.

However, **most of these quality improvements happen in the higher layers of the Internet, not the network layer.** The network layer remains mostly the same.

Security

Another thing to think about in network communication is safety. When the Internet first started, it wasn't a big deal because only a few people at universities used it for research, and outsiders couldn't access it. So, they didn't worry much about security in the network layer.

But now, security is a big deal. **To make the network layer more secure, they've created something called IPSec**, which adds a layer of protection to the network.

Packet switching how data moves around in a network. Think of it

like a switch that connects things, like how an electrical switch connects power. In

networking, there are two ways to do this: **circuit switching and packet**

switching. We use packet switching in the network layer because data is sent in packets here. Circuit

switching is used at a lower level. In the network layer, a

message is split into packets and sent through the network. The sender sends them one by one, and the receiver gets them one by one. **The receiver waits for**

all the packets from the same message before passing it to the next level. In a

packet-switched network, devices still need to figure out how to send the packets to the right place. There are

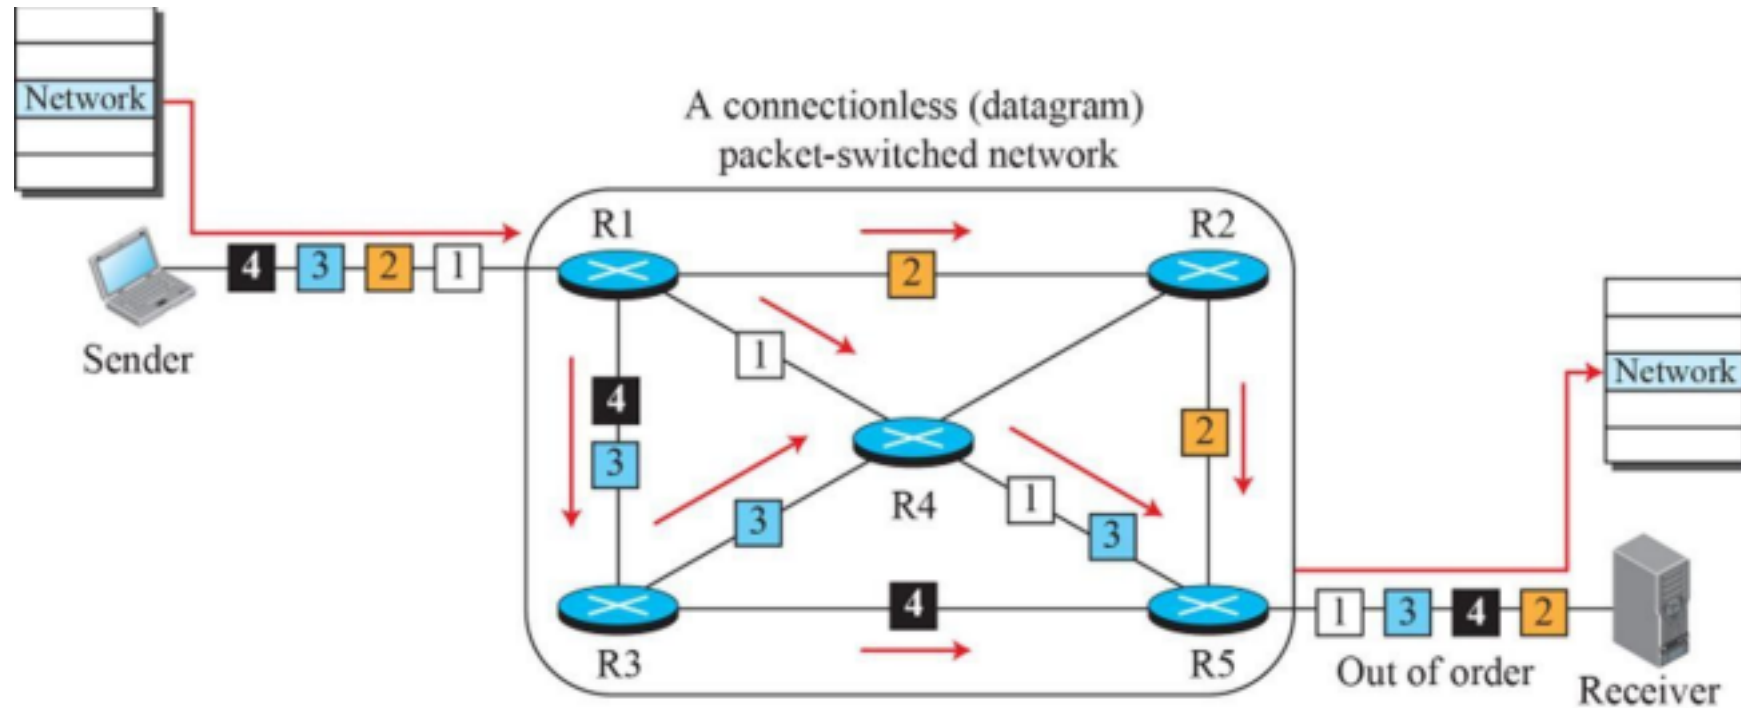
two ways to do this: • the datagram way

- the virtual circuit way

Datagram Approach: Connectionless Service

- When the Internet first started, the network layer was designed to keep things simple. It offered a service where each packet of data was treated separately, with no connection or relationship between them. **The network layer's job was just to get packets from one place to another**, and they could take different paths.
- Imagine each packet as its own independent traveler in the Internet journey. **They don't know each other, and they don't have to follow the same route.** Routers, which are like traffic cops for Internet data, help these packets find their way. They decide where to send each packet based on the destination address written on it.
- **The source address is there to help send an error message back to the sender if something goes wrong.** So, it's like each packet has a destination it's trying to reach, and routers help them get there. Think of it like sending

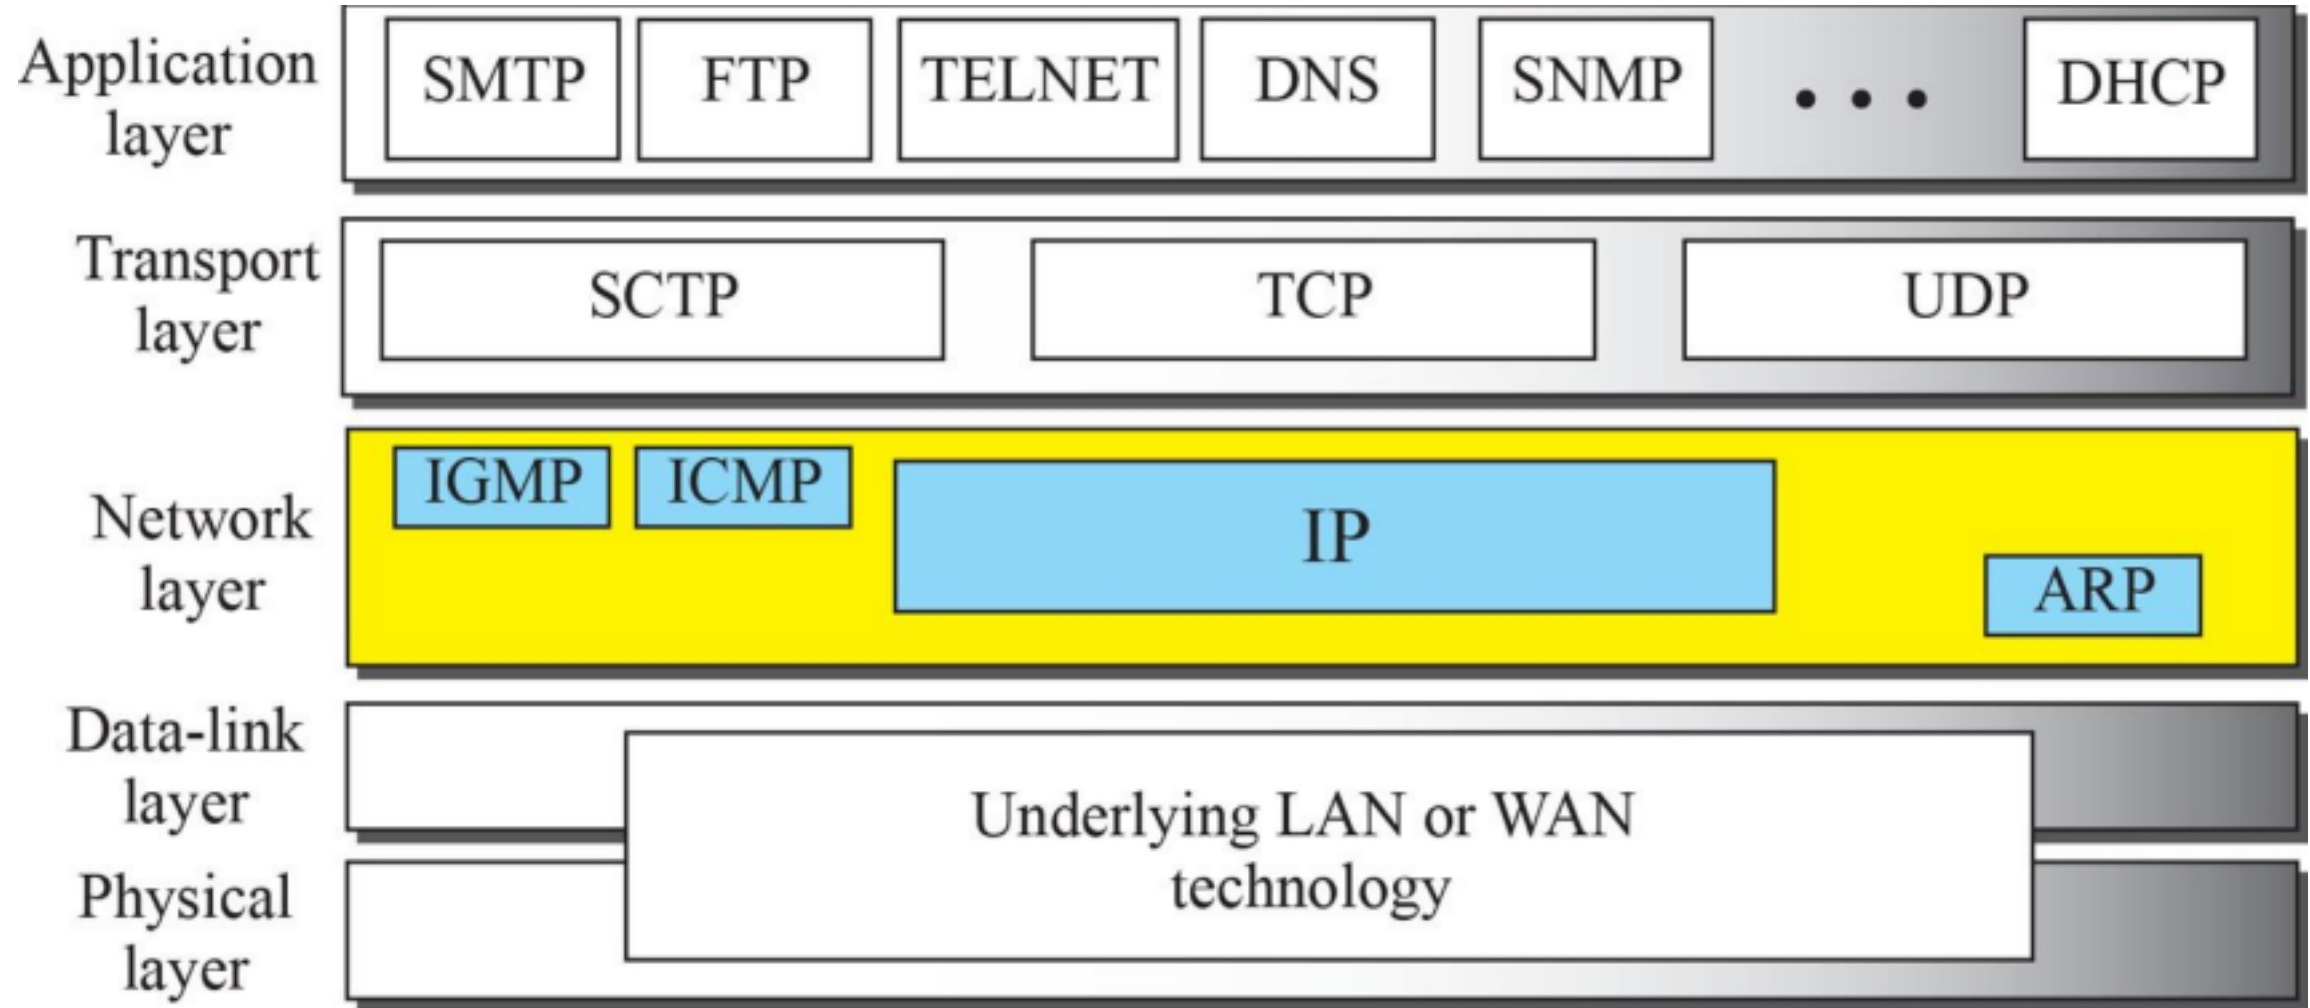
letters to different places



NETWORK-LAYER PROTOCOLS


- In this section, we show how the network layer is implemented in the TCP/IP protocol suite. The protocols in the network

layer have gone through several versions; in this section, we concentrate on the current version (4), in the last section of this chapter, we briefly discuss version 6, which is on the horizon.



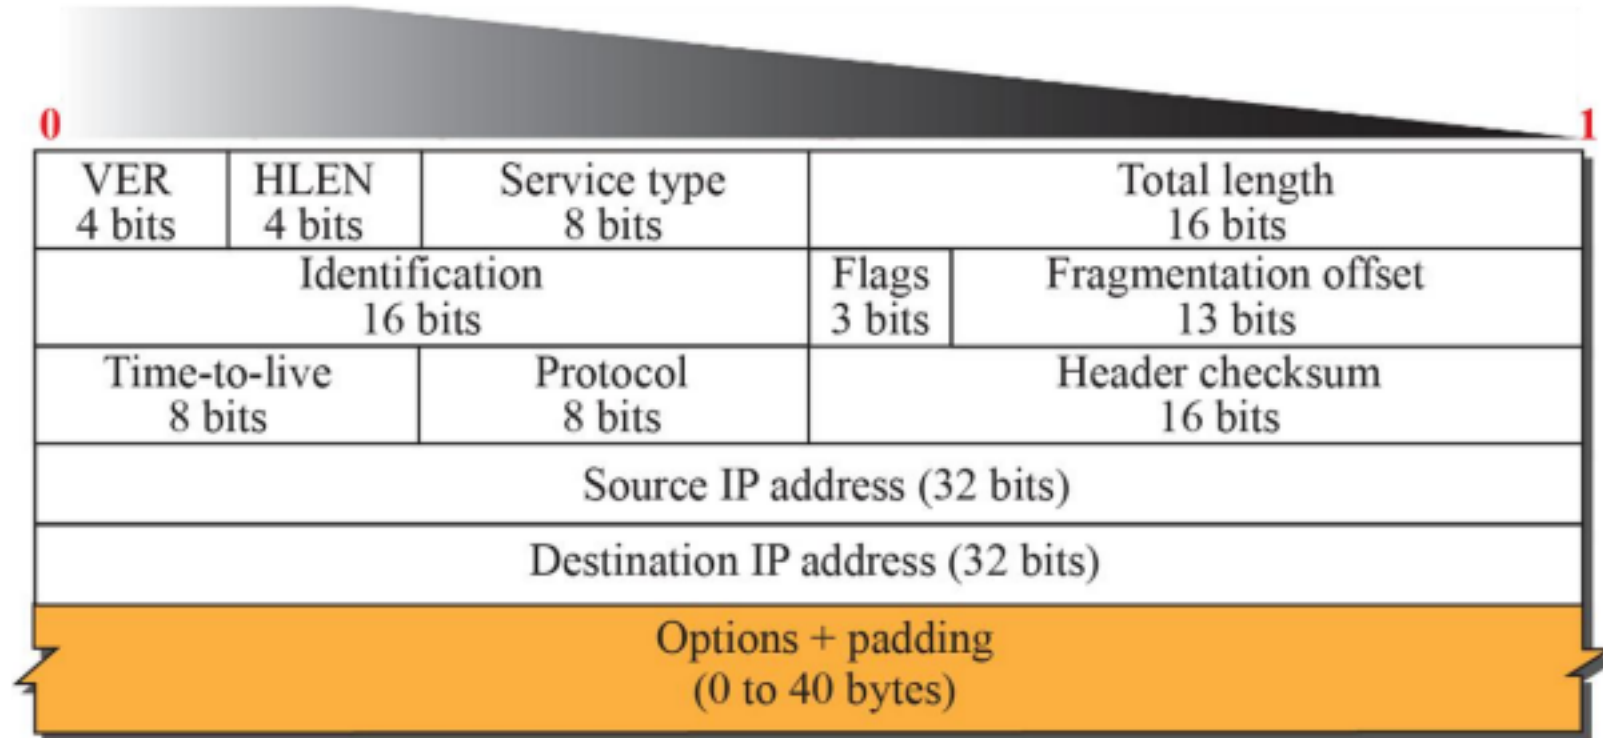
IPv4 DatagramFormat

- the IPv4 datagram format. A datagram is a variable-length packet consisting of two parts: **header and payload (data)**. The header is 20 to 60 bytes in length and contains information essential for routing and delivery. It is customary in TCP/IP to show the header in 4-byte sections



Legend

VER: version number
HLEN: header length
byte: 8 bits



b. Header format

• Version Number

The 4-bit version number (VER) field defines the version of

the IPv4 protocol, which, obviously, has the value of 4. • **Header Length (HLEN)**

In IPv4 is a 4-bit field that tells us how long the header of a datagram is in terms of 4-byte words. **Since the header's length can vary, this field helps devices figure out where the header ends and the actual data begins.** To fit this length into 4 bits, we calculate it in 4-byte words. To get the actual length in bytes, you need to multiply the value in this field by 4.

- **Service Type**

Used to be called Type of Service (TOS) in the IP header, and it determined how to treat the datagram. In the late 1990s, the IETF changed it to provide Differentiated Services (DiffServ). (RFC 2474)

- **Total Length:**

This 16-bit field tells us the size of the entire IP datagram in bytes, including both the

header and data. It can go up to 65,535 bytes, but most datagrams are smaller. It helps the receiving device know when the entire packet has arrived. • **Identification,**

Flags, and Fragmentation Offset: These three fields are connected to breaking up large IP datagrams into smaller pieces when the network can't handle their size.

We'll explain these fields in more detail in the next section when we talk about fragmentation. • **Time-to-live:** To prevent data packets from endlessly

traveling through the internet due to routing issues, the Time-to-live (TTL) field is used. It limits the number of routers a packet can pass through. The source sets an initial value, typically twice the maximum number of hops between hosts. Each router reduces this value by one. If it reaches zero, the router

drop the packet.

- The "**Protocol**" field in TCP/IP identifies which type of data is carried in a packet. It uses an 8-bit number assigned to each protocol, helping devices know how to handle the packet's content. This field works like a label, ensuring the right protocol deals with the packet, similar to how port numbers operate at the transport layer.
- **Header Checksum** in IP ensures the accuracy of the header but not the payload. Errors in the IP header can lead to misdirection or fragmentation issues. This 16-bit checksum is recalculated at each router to account for changes in certain fields like TTL. Checksum calculation is discussed in detail in Chapter 5 for error detection.
 - **Source and Destination Addresses:** These 32-bit fields in an IP header specify the source and destination IP addresses. The source host knows its IP address, while the

destination address is provided by the protocol using IP or through DNS. These addresses must stay unchanged during the datagram's journey. We'll dive deeper into IP addresses later in this chapter.

- **Options:** An IP datagram header can hold up to 40 bytes of optional information, useful for network testing and debugging. While not mandatory, all IP implementations must handle options if present. Having options can affect datagram handling and may require routers to recalculate the header checksum. We discuss one-byte and multi-byte options in more detail on the book's website.
- **Payload:** The payload, or data, is the primary content of a datagram. It's the information sent by other protocols using IP. Think of a datagram like a package: the payload is the package's content, and the header is

like the label on the package.

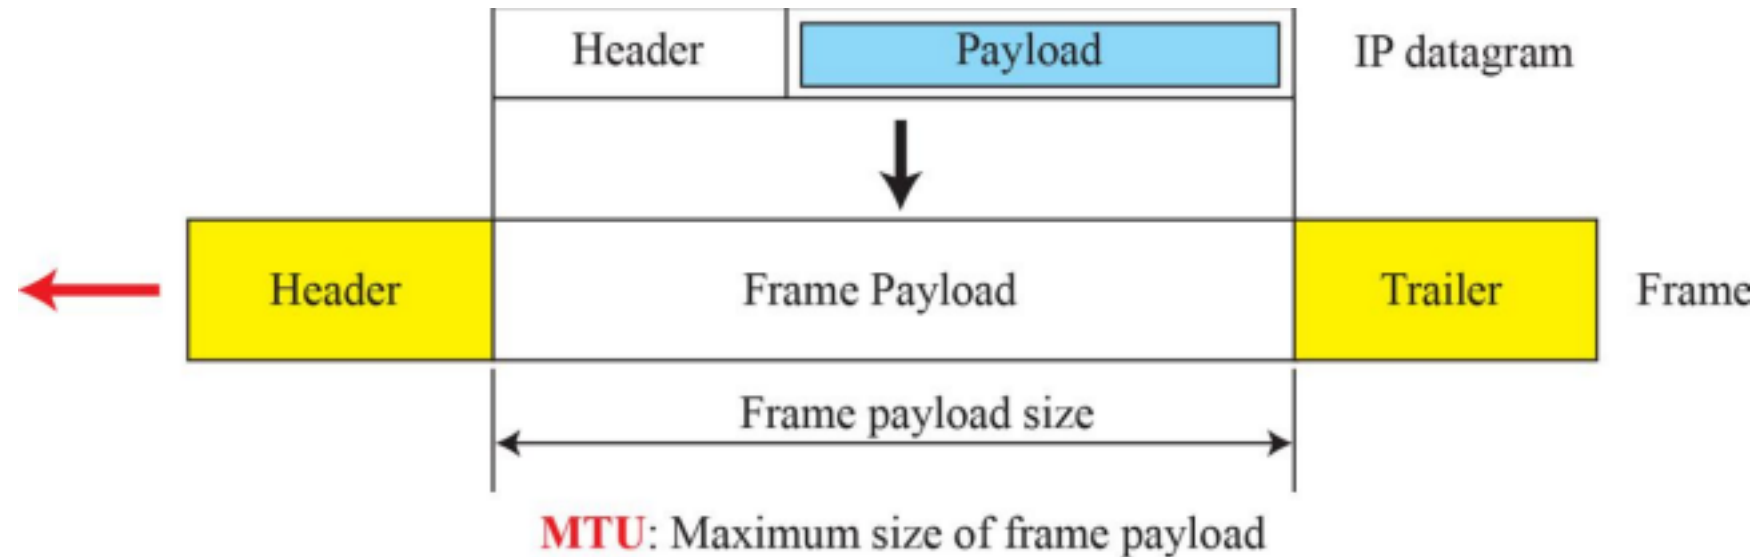
Fragmentation

- A datagram can move through various networks. When a router gets a datagram, it takes out the IP data and works on it. Then, it puts it into a new frame. The frame's look and size are different based on the network it's on. If a router links a local network (LAN) to a wide-area network (WAN), it gets a LAN-style frame and sends out a WAN-style frame.

Maximum Transfer Unit (MTU)

- Different networks use their own rules for packaging data. Each network type has a limit on how big a piece of data, called a payload, can be put into its package. This limit is called the

Maximum Transmission Unit (MTU), and it varies depending on the type of network. For example, a local network can handle a payload of up to 1500 bytes, but a wide-area network might allow more or less.



- The Internet Protocol (IP) allows data packets to be as large as 65,535 bytes, which is good for future networks. However, in current networks with smaller limits (MTUs), these large packets must be broken into smaller parts, called "fragmentation." If a network

with an even smaller limit is encountered, further splitting may occur. This means one data packet can be split multiple times before reaching its destination. • This splitting can be done by the source computer or any router along the way. However, the task of reassembling these pieces into the original packet is done only by the destination computer. Each piece acts as an independent packet as it travels through the network.

- Even though fragmented packets can take different paths, they are designed to eventually reach the destination. It's more efficient to reassemble them at the end, rather than during the journey, to avoid unnecessary complexity.
- Fragmentation mainly involves breaking up the payload of an IP data packet. Most header information, except for a few options, must be copied in each fragment. The device doing the fragmentation needs to adjust three fields:

flags, fragmentation offset, and total length. Everything else in the header remains the same. Also, the checksum value must be recalculated, regardless of how many times the data is split.

Fields Related to Fragmentation

- **Identification field**

The 16-bit ID field in a datagram helps identify where it comes from. To ensure each datagram is unique, the IP protocol uses a counter. This counter starts at a positive number and is increased by one for each datagram sent. As long as this counter stays in the computer's memory, uniqueness is guaranteed. When

a datagram gets split into smaller parts (fragments), all of them get the same ID number as the original datagram. This ID number helps the destination computer put the fragments back together. It knows that all fragments with the same ID should be assembled into one datagram.

- **Flag field**

The 3-bit flags field has three flags. The first one isn't used. The second flag, called the "**do not fragment**" bit (D bit), when set to 1, means the datagram shouldn't be broken into pieces. If it can't fit through a network, it's discarded, and an error message is sent back. When the D bit is set to 0, the datagram can be split into smaller parts if needed.

The third flag, called the "**more fragment**" bit (M bit), is set to 1 when there are more fragments to come after this one. If it's set to 0, it means this is

the last part or the only part of the datagram.

Offset

The 13-bit **fragmentation offset field** tells us where this piece fits in the whole **dataragram**, counting in chunks of 8 bytes. Imagine we have a datagram with 4000 bytes split into three parts. Here's how it works:

- The first piece covers bytes 0 to 1399 and has an offset of 0.
- The second piece covers bytes 1400 to 2799 and has an offset of 175.
- The third piece covers bytes 2800 to 3999 and has an offset of 350.

Imagine the fragments are like pieces of a puzzle sent from one place to another. Even if they arrive in a different order, the receiving end can still put the puzzle together correctly. This happens because:

- The identification and flags fields are the same in all fragments, except for the last

- The offset field in each fragment shows its position relative to the original data. If a fragment itself gets split into smaller pieces, the offset is still relative to the original data. So, even if pieces take different paths and arrive out of order, they can be correctly reassembled. **To reassemble, follow these steps:**
 - a. The first fragment has an offset of 0.
 - b. Divide the length of the first fragment by 8 to get the offset of the second fragment.
 - c. Divide the total length of the first and second fragments by 8 to get the offset of the third fragment.
 - d. Continue this process. The last fragment has its "more" bit set to 0.
- Remember that the value of the offset is measured in units of 8 bytes. **This is**

done because the length of the offset field is only 13 bits long and cannot represent a sequence of bytes greater than 8191. This forces hosts or routers that fragment datagrams to choose the size of each fragment so that the first byte number is divisible by 8.

Q - A datagram of 3000 B (20B of IP header + 2980B IP payload) reached at router and must be forward to link with MTU of 500B. How many fragments will be generated and also write: MF, Offset, total length value for all.

Security of IPv4 Datagrams

The IPv4 protocol, as well as the whole Internet, was started when the Internet users trusted each other. No security was provided for the IPv4 protocol. Today, however, the situation is different; the Internet is not secure anymore. There are three security issues that are particularly applicable to the IP protocol:

- packet sniffing

- packet modification
- IP spoofing.

Packet Sniffing

- Packet Sniffing **An intruder may intercept an IP packet and make a copy of it.** Packet sniffing is a passive attack, in which the attacker **does not change the contents of the packet.** This type of attack is very difficult to detect because the sender and the receiver may never know that the packet has been copied. Although packet sniffing cannot be stopped, **encryption of the packet can make the attacker's effort useless.** The attacker may still sniff the packet, but the content is not detectable

Packet Modification

- The 2nd type of attack is to modify the packet. The **attacker intercepts the packet, changes its contents, and sends the new packet to the receiver**. The receiver believes that the packet is coming from the original sender. This type of attack can be detected using a data integrity mechanism. The receiver, before opening and using the contents of the message, can use this mechanism to make sure that the packet has not been changed during the transmission.

IP Spoofing

- A person with malicious intent can pretend to be someone else and send an **IP packet that looks like it's from a different computer**. For example, they can send a packet to a bank, making it seem like it's from one of the bank's customers.

IPSec

The IP packets today can be protected from the previously mentioned attacks using a protocol called IPSec (IP Security). This protocol, which is used in conjunction with the IP protocol, creates a connection-oriented service between two entities in which they can exchange IP packets without worrying about the three attacks

discussed above. • **Defining Algorithms and Keys.** The two entities that want

to create a secure channel between themselves can agree on some

available algorithms and keys to be used for security purposes. • **Packet Encryption.**

The packets exchanged between two parties can be encrypted for privacy using one

of the encryption algorithms and a shared key agreed upon in the first step. This makes the packet sniffing attack useless.

- **Data Integrity.** Data integrity guarantees that the packet is not modified during the transmission. If the received packet does not pass the data integrity test, it is discarded. This prevents the second attack, packet modification, described above.
- **Origin Authentication.** IPSec can authenticate the origin of the packet to be sure that the packet is not created by an imposter. This can prevent IP spoofing attacks as described above.

IPv4 Addresses

- The Internet address, known as an IP address, is like a label for devices on the Internet. It's a 32-bit code that tells the Internet where a device is. It's

important to note **that this address is for the connection, not the device itself. If a device switches networks, its IP address can change.** Each IP address is unique and corresponds to one Internet connection. So, **if a device has two Internet connections, it has two IP addresses.** These addresses are universal, meaning any Internet-connected device must recognize them.

Address Space

- A protocol like IPv4 that defines addresses has an address space. An address space is the total number of addresses used by the protocol. If a protocol uses b bits to define an address, the address space is 2^b because each bit can have two different values (0 or 1). IPv4 uses 32-bit addresses, which means that the address space is $2^{(32)}$ or 4,294,967,296 (more than four billion). If there were no restrictions, more than 4 billion devices could be connected to the Internet.



Notation

IPv4 addresses can be written in three common ways: binary(0sand1s), dotted-decimal (numbers separated by dots, base 256), andhexadecimal (numbers and letters, base 16).

- In binary, it's 32 bits with spaces to make it readable usuallybetweeneach octet (8 bits).
- Dotted-decimal is the usual way, with dots separating numbersfrom0to 255.
- Hexadecimal uses numbers and letters and is commoninnetworkprogramming. Figure displays an IP address in these threeways.









Conversion D-B, B-D



Example 1



Example 2



Hierarchy in Addressing

In networks like telephone or postal systems, the addressing works in a structured way. For example, a postal address includes country, state, city, street, house number, and recipient's name. Similarly, a phone number has country code, area code, local exchange, and the specific connection.

Now, in the case of a 32-bit IPv4 address, it's also structured but in a simpler way. It's divided into two parts: the first part (prefix) defines the network, and the second part (suffix) defines the device on that network. The prefix can be either a fixed or variable length.

In the past, **IPv4 used fixed-length prefixes (classful addressing), but that's outdated. Nowadays, it uses variable-length prefixes (classless addressing).** We'll first touch on the old method and then focus on the new one.



Classful Addressing

- When the Internet started, an IPv4 address was designed with a fixed length prefix, but to accommodate both small and large networks, three fixed-length prefixes were designed instead of one ($n=8$, $n=16$, and $n=24$). The whole address space was divided into five classes (class A, B, C, D, and E), as shown in Figure. This scheme is referred to as classful addressing.



- **Class A**

the network length is 8 bits, but since the first bit, which is 0, defines the class, we can have only seven bits as the network identifier. This means there are only $2^7 = 128$ networks in the world that can have a class A address. • **Class B**

the network length is 16 bits, but since the first two bits, which are (10)_2, define the class, we can have only 14 bits as the network identifier. This means there are only $2^{14} = 16,384$ networks in the world that can have a class B address. • **Class C**

All addresses that start with (110)_2 belong to class C. In class C, the network length is 24 bits, but since three bits define the class, we can have only 21 bits as the network identifier. This means there are $2^{21} = 2,097,152$ networks in the world that can have a class C address.

• **Class D and Class E**

Class D is not divided into prefix and suffix. It is used for multicast addresses. All addresses that start with 1111 in binary belong to class E. As in Class D, Class E is not divided into prefix and suffix and is used as a reserve.







Address Depletion

- Classful addressing became outdated because it led to address shortage. Internet addresses were not handed out efficiently, so we ran out of them fast. For instance, class A addresses were given to only 128 organizations, but each one got a massive chunk of addresses, which most didn't need. Class B addresses were for midsize organizations, but many went unused. Class C addresses had too few usable addresses for companies. Class E addresses were hardly used at all, wasting the entire class. This inefficiency in address distribution caused problems, which is why we moved to a different system.

Subnetting and Supernetting

- To deal with the problem of running out of addresses, two strategies were suggested: subnetting and supernetting.
- Subnetting involves breaking up a big address block (like class A or class B) into smaller parts, creating more subnetworks with shorter prefixes. This also allowed unused addresses to be shared among different organizations. However, many large organizations didn't like sharing their unused addresses, so this idea didn't work well.
- Supernetting, on the other hand, aimed to combine multiple class C blocks into a larger one to provide more addresses

to organizations. But this made it harder to route data packets effectively, so it wasn't a successful solution either.

Advantage of Classful Addressing

- Classful addressing, despite its issues and eventual obsolescence, had one clear benefit:

it made it easy to determine the class of an address and, because each class had a fixed prefix length, you could instantly know the prefix length from the address itself. In simple terms, the address itself contained all the information needed to figure out its prefix and suffix; no extra details were required.

Classless Addressing

Address Deletion and the need for change

- Subnetting and supernetting in classful addressing didn't effectively address address depletion.
- The growth of the Internet highlighted the need for a larger address space.
- IPv6 was developed as a long-term solution, but a short-term solution called classless addressing was also introduced.
- Classless addressing removed the class distinctions in address distribution to address depletion.
- Internet Service Providers (ISPs) played a role in the adoption of classless addressing.

Classless Addressing: A flexible solution

- Classless addressing introduced variable-length address blocks without class distinctions.

- Address space is divided into variable-length blocks, where the prefix defines the network, and the suffix defines the device. • Blocks can vary in size from 2^0 to 2^{32} addresses.
- Prefix length ranges from 0 to 32, with smaller prefixes indicating larger networks and vice versa.
- Classless addressing can be applied to classful addressing, making classful addressing a special case of classless addressing.




Prefix Length: Slash Notation • In classless addressing,

determining the prefix length is crucial.

- Unlike classful addressing, the prefix length is not inherent in the address.
- To specify the prefix length, it's added to the address using slash notation or CIDR (pronounced cider) strategy.
- The format is like this: address/prefix_length.
 - This means that in classless addressing, an address alone doesn't fully define its network; we also need to provide the prefix

length.



Extracting Information from an Address

Given any address in the block, we normally like to know three pieces of information about the block to which the address belongs: • the number of addresses

- the first address in the block
- the last address.

If the value of prefix length, n , is given, we can easily find these three pieces of information



Method for extracting
information from an address

- The number of addresses in the block is found as $N = 2^{(32-n)}$.
- To find the first address, we keep the n leftmost bits and set the $(32-n)$ rightmost bits all to 0s.
- To find the last address, we keep the n leftmost bits and set the $(32-n)$ rightmost bits all to 1s.

Example

A classless address is given as 167.199.170.82/27. We can find the above three pieces of information as follows. The number of addresses in the network is $2^{32-n} = 2^5 = 32$ addresses. The first address can be found by keeping the first 27 bits and changing the rest of the bits to 0s.



The last address can be found by keeping the first 27 bits and

changing the rest of the bits to 1s.



Address Mask

Another way to find the first and last addresses in the block is to use the address mask. The address mask is a 32-bit number in which the leftmost bits are set to 1s and the rest of the bits ($32 - n$) are set to 0s. A computer can easily find the address mask because it is the complement of $(2^{(32-n)} - 1)$. The reason for defining a mask in this way is that it can be used by a computer program to extract the information in a block, using the three bit-wise operations NOT, AND, and OR.

1. The number of addresses in the block $N = \text{NOT}(\text{Mask}) + 1$.
2. The first address in the block = (Any address in the block) AND (Mask).
3. The last address in the block = (Any address in the block) OR [(NOT (Mask))].

Example

The mask in dotted-decimal notation is 255.255.255.224 The AND, OR, and NOT operations can be applied to individual bytes using calculators.

Number of addresses in the block: $N = \text{NOT}(\text{mask}) + 1 = 0.0.0.31 + 1 = 32$ addresses

First address: First = (address) AND (mask) = 167.199.170. 82
Last address: Last = (address) OR (NOT mask) = 167.199.170. 255

Network Address

- The provided examples demonstrate our ability to retrieve block-related information from any given address.
- The network address, which is the first address, is vital for

routing packets to their destination networks.

- Let's assume there are m networks in an internet, each connected to a router with m interfaces.
- When a packet arrives at the router from any source host, the router must determine which network to send the packet to and which interface to use.
- The packet employs a different strategy to reach its final destination host once it reaches the network.
- Figure visually represents this process.
- **After identifying the network address, the router checks its forwarding table to find the correct interface for sending out the packet.**
- Each network is uniquely identified by its network address.



Block Allocation

The next concern with classless addressing is how blocks are handed out. The job of

giving out these blocks falls on a global authority named ICANN. But, ICANN doesn't usually hand out addresses to regular internet users. Instead, it gives a big chunk of addresses to an ISP (or a similar big organization that acts like an ISP). To make CIDR work well, there are two important rules that need to be followed for these allocated blocks.

1. The number of requested addresses, N , needs to be a power of 2. The reason is that $N = 2^{(32 - n)}$ or $n = 32 - \log_2(N)$. If N is not a power of 2, we cannot have an integer value for n .
2. The requested block needs to be allocated where there are a contiguous number of available addresses in the address space. However, there is a restriction on choosing the first address in the block. The first address needs to be divisible by the number of addresses in the block. The reason is that the first address needs to be the prefix followed by $(32 - n)$ number of 0s. The decimal value of the first address is then

$$\text{first address} = (\text{prefix in decimal}) \times 2^{(32 - n)} = (\text{prefix in decimal}) \times N$$

Example

An ISP has requested a block of 1000 addresses.

Since 1000 is not a power of 2, 1024 addresses are granted. The prefix length is calculated as $n = 32 - \log_2(1024) = 22$. An available block, 18.14.12.0/22, is granted to the ISP. It can be seen that the first address in decimal is 302,910,464, which is divisible by 1024.

Subnetting

Subnetting allows for the creation of additional layers in the network hierarchy. When an organization or ISP is given a block of addresses, **they can split it into smaller sub-blocks and allocate each of these to separate subnetworks.**

Importantly, there's no limit to how deep this hierarchy can go. A subnetwork can be further divided into sub-subnetworks, and the process can continue with sub-sub-subnetworks, and so forth.



Finding Information about EachSubnetwork

After designing the subnetworks, the information about each subnetwork, such as first and last address, can be found using the process we described to find the information about each network in the Internet.

EXAMPLE: An organization is granted a block of addresses with the beginning address 14.24.74.0/24. The organization needs to have 3 subblocks of addresses to use in its three subnets: one subblock of 10 addresses, one subblock of 60 addresses, and one subblock of 120 addresses. Design the subblocks.

Solution

There are 2

$2^{32-24} = 256$ addresses in this block. The first address is 14.24.74.0/24; the last address is 14.24.74.255/24. To satisfy the third requirement, we assign addresses to subblocks, starting with the

largest and ending with the smallest one.

a. The number of addresses in the largest subblock, which requires 120 addresses, is not a power of 2. We allocate 128 addresses. The subnet mask for this subnet can be found as $n_1 = 32 - \log_2 128 = 25$. The first address in this block is 14.24.74.0/25; the last address is 14.24.74.127/25.

b. The number of addresses in the second largest subblock, which requires 60 addresses, is not a power of 2 either. We allocate 64 addresses. The subnet mask for this subnet can be found as $n_2 = 32 - \log_2 64 = 26$. The first address in this block is 14.24.74.128/26; the last address is 14.24.74.191/26.

c. The number of addresses in the smallest subblock, which requires 10 addresses, is not a power of 2. We allocate 16 addresses. The subnet mask for this subnet can be found as $n_3 = 32 - \log_2 16 = 28$. The first address in this block is 14.24.74.207/28; the last address is

14.24.74.207/28.

If we add all addresses in the previous subblocks, the result is 208 addresses, which means 48 addresses are left in reserve. The first address in this range is 14.24.74.208. The last address is 14.24.74.255. We don't know about the prefix length yet. Figure 4.36 shows the configuration of blocks. We have shown the first address in each block.